

1 Image classification

NVIDIA System Management Interface ซึ่งใช้แสดงข้อมูลและจัดการการใช้งานของ GPU บนระบบ

```
! nvidia-smi
```

โค้ดนี้ทำหน้าที่ปิดการแจ้งเตือน (warnings) ที่เกิดจากไลบรารี scikit-learn เมื่อมีการคำนวณเมตริกที่ไม่ได้กำหนดไว้

```
from sklearn.exceptions import UndefinedMetricWarning

def warn(*args, **kwargs):
    pass

import warnings
warnings.warn = warn
```

โค้ดนี้ใช้ไลบรารี PyTorch และ torchvision ที่เป็นไลบรารีที่ถูกพัฒนาขึ้นสำหรับงานด้าน deep learning และ computer vision ใน Python

```
import torch
import torchvision
import torchvision.transforms as transforms
```

ตรวจสอบว่าเครื่องคอมพิวเตอร์ที่รันโค้ดนี้มี GPU (กราฟิกประมวลผลหลัก) ที่สามารถใช้งานได้หรือไม่ โดยทำการตั้งค่าตัวแปร `device` ให้กับ CUDA GPU ถ้ามี, และถ้าไม่มีก็จะให้ใช้ CPU แทน

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

# Assuming that we are on a CUDA machine, this should print a CUDA device:

print(device)
```

แปลงรูปภาพเป็น tensor และปรับขนาดของรูปภาพเป็นขนาดที่กำหนด

```
transform = transforms.Compose( # transform is from torchvision (only for image)
    [transforms.ToTensor(), # image to tensor --> divide by 255
     transforms.Resize((32, 32))])

batch_size = 32
```

สร้าง Dataset สำหรับการฝึกและ Validation และ สร้าง DataLoader สำหรับ Trainset และ Validation

torchvision.datasets

```
<module 'torchvision.datasets' from '/usr/local/lib/python3.10/dist-packages/torchvision/datasets/__init__.py'>
```

```
trainvalset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
trainset, valset = torch.utils.data.random_split(trainvalset, [40000, 10000])
```

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True)
valloader = torch.utils.data.DataLoader(valset, batch_size=batch_size, shuffle=False)
```

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False)
```

```
#classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

ใช้เพื่อดูขนาด (จำนวนตัวอย่าง) ของ Dataset สำหรับการฝึก (trainset), Dataset สำหรับ Validation และ Dataset สำหรับการทดสอบ (testset)

```
trainset.__len__(), valset.__len__(), testset.__len__()
```

```
(40000, 10000, 10000)
```

def imshow(img): เป็นการแสดงรูปภาพ ส่วน iter(trainloader) คือการสุ่มรูปภาพมาจาก trainloader โดยกำหนดให้ 1 แถวแสดง 9 ช่อง

```
import matplotlib.pyplot as plt
import numpy as np

# functions to show an image
def imshow(img):
    npimg = img.numpy()
    plt.figure(figsize=(16,16))
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataliter = iter(trainloader)
images, labels = next(dataliter)

# show images
nrow = 9
imshow(torchvision.utils.make_grid(images, nrow = nrow))
```



พิมพ์ Label ของรูปภาพที่ถูกดึงมาจาก `trainloader` ในรูปแบบของกริด โดยใช้ลูปเพื่อแสดง Label ของแต่ละรูปภาพในแต่ละแถวของกริด

```
# print labels
for i in range(batch_size//nrow + 1 if batch_size % nrow else 0):
    print(' '.join(f'{labels[i*nrow+j]:<3}' for j in range(min(batch_size - i*nrow, nrow))))
```

```
0  9  7  8  6  8  4  3  1
0  6  8  5  6  2  2  9  5
4  6  1  4  4  9  3  7  3
0  7  6  2  4
```

สร้างและกำหนดโมเดล CNN ที่ใช้ในการจำแนกประเภทของภาพ (image classification) ด้วย PyTorch และ มีลักษณะของชั้น convolution, pooling, และ fully connected layers

```
import torch.nn as nn
import torch.nn.functional as F

class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5) # 3 input channels, 6 output channels, 5*5 kernel size
        self.pool = nn.MaxPool2d(2, 2) # 2*2 kernel size, 2 strides
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(400, 120) # dense input 400 (16*5), output 120

        self.fc2 = nn.Linear(120, 84) # dense input 120, output 84
        self.fc3 = nn.Linear(84, 10) # dense input 84, output 10
        self.softmax = torch.nn.Softmax(dim=1) # perform softmax at dim[1] (batch,class)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, start_dim=1) # flatten all dimensions (dim[1]) except batch (dim[0])
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        x = self.softmax(x)
        return x

net = CNN().to(device)
```

ติดตั้งไลบรารี torchinfo ที่ใช้สำหรับแสดงข้อมูลเกี่ยวกับโมเดล PyTorch

```
!pip install torchinfo
```

ใช้ torchinfo เพื่อแสดงข้อมูลสรุปเกี่ยวกับโมเดล PyTorch

```
from torchinfo import summary as summary_info
print(summary_info(net, input_size = (32, 3, 32, 32))) # (batchsize,channel,width,height)
net = net.to(device)
```

```
=====
Layer (type:depth-idx)                   Output Shape          Param #
-----
CNN
|---Conv2d: 1-1                          [32, 6, 28, 28]       456
|---MaxPool2d: 1-2                       [32, 6, 14, 14]       --
|---Conv2d: 1-3                          [32, 16, 10, 10]      2,416
|---MaxPool2d: 1-4                       [32, 16, 5, 5]        --
|---Linear: 1-5                           [32, 120]             48,120
|---Linear: 1-6                           [32, 84]              10,164
|---Linear: 1-7                           [32, 10]              850
|---Softmax: 1-8                         [32, 10]              --
=====
Total params: 62,006
Trainable params: 62,006
Non-trainable params: 0
Total mult-adds (M): 21.06
=====
Input size (MB): 0.39
Forward/backward pass size (MB): 1.67
Params size (MB): 0.25
Estimated Total Size (MB): 2.31
=====
```

ฝึกโมเดล PyTorch (criterion), และใช้ตัวอัปเดตพารามิเตอร์ของโมเดลด้วยวิธีการ Stochastic Gradient Descent (SGD) ที่กำหนด learning rate เป็น 1e-2 และ momentum เป็น 0.9

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=1e-2, momentum=0.9)
```

ฝึกโมเดล CNN สำหรับการจำแนกภาพด้วย PyTorch โดยใช้ SGD ในแต่ละ epoch และบันทึกผลลัพธ์การฝึกและการทดสอบ เพื่อเลือกโมเดลที่มีค่า validation loss ต่ำที่สุด

```
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from tqdm.notebook import tqdm

epochs = 20

history_train = {'loss':np.zeros(epochs), 'acc':np.zeros(epochs), 'f1-score':np.zeros(epochs)}
history_val = {'loss':np.zeros(epochs), 'acc':np.zeros(epochs), 'f1-score':np.zeros(epochs)}
min_val_loss = 1e10
PATH = './CNN_CIFAR10.pth'

for epoch in range(epochs): # loop over the dataset multiple times

    print(f'epoch {epoch + 1} \nTraining ...')
    y_predict = list()
    y_labels = list()
    training_loss = 0.0
    n = 0
    net.train()
    for data in tqdm(trainloader):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs) # forward
        loss = criterion(outputs, labels) # calculate loss from forward pass
        loss.backward() # just calculate
        optimizer.step() # update weights here

    # aggregate statistics
    training_loss += loss.item()
    n+=1

    y_labels += list(labels.cpu().numpy())
    y_predict += list(outputs.argmax(dim=1).cpu().numpy())

    # print statistics
    report = classification_report(y_labels, y_predict, digits = 4, output_dict = True)
```

```
acc = report["accuracy"]
f1 = report["weighted avg"]["f1-score"]
support = report["weighted avg"]["support"]
training_loss /= n
print(f'training loss: (training_loss:.4), acc: (acc*100:.4)%, f1-score: (f1*100:.4)%, support: {support}')
history_train['loss'][epoch] = training_loss
history_train['acc'][epoch] = acc
history_train['f1-score'][epoch] = f1

print('validating ...')
net.eval()
y_predict = list()
y_labels = list()
validation_loss = 0.0
n = 0
with torch.no_grad():
    for data in tqdm(valloader):
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)

        outputs = net(inputs)
        loss = criterion(outputs, labels)
        validation_loss += loss.item()

        y_labels += list(labels.cpu().numpy())
        y_predict += list(outputs.argmax(dim=1).cpu().numpy())
        n+=1

    # print statistics
    report = classification_report(y_labels, y_predict, digits = 4, output_dict = True)
    acc = report["accuracy"]
    f1 = report["weighted avg"]["f1-score"]
    support = report["weighted avg"]["support"]
    validation_loss /= n
    print(f'validation loss: (validation_loss:.4), acc: (acc*100:.4)%, f1-score: (f1*100:.4)%, support: {support}')
    history_val['loss'][epoch] = validation_loss
    history_val['acc'][epoch] = acc
    history_val['f1-score'][epoch] = f1

    # save min validation loss
    if validation_loss < min_val_loss:
        torch.save(net.state_dict(), PATH)
        min_val_loss = validation_loss

print('Finished Training')
```

```

epoch 1
Training ...
100% ██████████ 1250/1250 [00:17:00.00, 57.72u/s]
training loss: 2.293, acc: 11.38%, f1-score: 6.614%, support: 40000
validating ...
100% ██████████ 313/313 [00:04:00.00, 55.85u/s]
validation loss: 2.259, acc: 16.6%, f1-score: 7.761%, support: 10000
epoch 2
Training ...
100% ██████████ 1250/1250 [00:12:00.00, 130.47u/s]
training loss: 2.209, acc: 23.26%, f1-score: 19.79%, support: 40000
validating ...
100% ██████████ 313/313 [00:02:00.00, 98.95u/s]
validation loss: 2.169, acc: 28.2%, f1-score: 25.62%, support: 10000
epoch 3
Training ...
100% ██████████ 1250/1250 [00:09:00.00, 131.78u/s]
training loss: 2.152, acc: 29.94%, f1-score: 27.62%, support: 40000
validating ...
100% ██████████ 313/313 [00:02:00.00, 149.60u/s]
validation loss: 2.146, acc: 30.77%, f1-score: 28.81%, support: 10000
epoch 4
Training ...
100% ██████████ 1250/1250 [00:09:00.00, 110.40u/s]
training loss: 2.111, acc: 34.21%, f1-score: 31.96%, support: 40000
validating ...
100% ██████████ 313/313 [00:02:00.00, 160.95u/s]
validation loss: 2.107, acc: 34.31%, f1-score: 32.38%, support: 10000
epoch 5
Training ...
100% ██████████ 1250/1250 [00:09:00.00, 99.59u/s]
training loss: 2.085, acc: 36.93%, f1-score: 34.8%, support: 40000
validating ...

```

```

100% ██████████ 1250/1250 [00:10:00.00, 126.32u/s]
training loss: 1.944, acc: 51.3%, f1-score: 51.04%, support: 40000
validating ...
100% ██████████ 313/313 [00:01:00.00, 160.14u/s]
validation loss: 1.967, acc: 49.19%, f1-score: 48.95%, support: 10000
epoch 17
Training ...
100% ██████████ 1250/1250 [00:10:00.00, 126.68u/s]
training loss: 1.946, acc: 51.12%, f1-score: 50.81%, support: 40000
validating ...
100% ██████████ 313/313 [00:02:00.00, 163.97u/s]
validation loss: 1.979, acc: 48.1%, f1-score: 47.15%, support: 10000
epoch 18
Training ...
100% ██████████ 1250/1250 [00:10:00.00, 129.91u/s]
training loss: 1.952, acc: 50.65%, f1-score: 50.34%, support: 40000
validating ...
100% ██████████ 313/313 [00:01:00.00, 162.47u/s]
validation loss: 1.979, acc: 47.92%, f1-score: 47.09%, support: 10000
epoch 19
Training ...
100% ██████████ 1250/1250 [00:10:00.00, 130.24u/s]
training loss: 1.945, acc: 51.31%, f1-score: 50.93%, support: 40000
validating ...
100% ██████████ 313/313 [00:01:00.00, 164.15u/s]
validation loss: 1.95, acc: 50.83%, f1-score: 50.47%, support: 10000
epoch 20
Training ...
100% ██████████ 1250/1250 [00:10:00.00, 129.08u/s]
training loss: 1.956, acc: 50.22%, f1-score: 50.03%, support: 40000
validating ...
100% ██████████ 313/313 [00:01:00.00, 148.63u/s]
validation loss: 1.991, acc: 46.57%, f1-score: 45.03%, support: 10000
Finished Training

```

min_val_loss เก็บค่า validation loss ที่ต่ำที่สุดที่พบในระหว่างการฝึกโมเดล

```
min_val_loss
```


```
1.948860833058342
```

ใช้ Matplotlib เพื่อสร้างกราฟเส้นแสดงค่า loss, accuracy (acc), และ f1-score ของโมเดลในระหว่างการฝึกและการทดสอบ

```

fig, axs = plt.subplots(3, figsize=(6,10))
# loss
axs[0].plot(history_train['loss'], label='training')
axs[0].plot(history_val['loss'], label='validation')
axs[0].set_title("loss")
axs[0].legend()
# acc
axs[1].plot(history_train['acc'], label='training')
axs[1].plot(history_val['acc'], label='validation')
axs[1].set_title("acc")
axs[1].legend()
# f1-score
axs[2].plot(history_train['f1-score'], label='training')
axs[2].plot(history_val['f1-score'], label='validation')
axs[2].set_title("f1-score")
axs[2].legend()
plt.show()

```



สร้างโมเดล CNN และโหลดน้ำหนักที่ถูกบันทึกไว้ (PATH) เพื่อนำมาใช้ในการทดสอบ

```

net = CNN().to(device)
net.load_state_dict(torch.load(PATH))

```

```
<All keys matched successfully>
```

ทดสอบโมเดล CNN ด้วยชุดข้อมูลทดสอบและคำนวณค่า test loss รวมถึงแสดงผล confusion matrix และรายงานประสิทธิภาพด้วย classification report โดยใช้ sklearn

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

print('testing ...')
y_predict = list()
y_labels = list()
test_loss = 0.0
n = 0
with torch.no_grad():
    for data in tqdm(testloader):
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)

        outputs = net(inputs)
        loss = criterion(outputs, labels)
        test_loss += loss.item()

        y_labels += list(labels.cpu().numpy())
        y_predict += list(outputs.argmax(dim=1).cpu().numpy())
        n+=1


# print statistics
test_loss /= n
print(f"testing loss: {test_loss:.4}")

report = classification_report(y_labels, y_predict, digits = 4)
M = confusion_matrix(y_labels, y_predict)
print(report)
disp = ConfusionMatrixDisplay(confusion_matrix=M)
#acc = report["accuracy"]
#f1 = report["weighted avg"]["f1-score"]
#support = report["weighted avg"]["support"]
#test_loss /= n
#print(f"validation loss: {test_loss:.4}, acc: {acc*100:.4}%, f1-score: {f1*100:.4}%, support: {support}")
```

```
testing ...
100% ██████████ 313/313 [00:02<00:00, 145.34it/s]
testing loss: 1.946
```

	precision	recall	f1-score	support
0	0.5334	0.6140	0.5709	1000
1	0.5203	0.7680	0.6204	1000
2	0.3898	0.4260	0.4071	1000
3	0.3953	0.2680	0.3194	1000
4	0.5096	0.3190	0.3924	1000
5	0.4795	0.3390	0.3972	1000
6	0.5271	0.6230	0.5710	1000
7	0.5314	0.6520	0.5855	1000
8	0.5946	0.6630	0.6270	1000
9	0.5812	0.4330	0.4963	1000
accuracy			0.5105	10000
macro avg	0.5062	0.5105	0.4987	10000
weighted avg	0.5062	0.5105	0.4987	10000

ใช้ Matplotlib เพื่อแสดงผลกราฟของ confusion matrix ที่ถูกสร้างขึ้นและบันทึกใน `disp` จาก sklearn



2_Image_classification_Animal_EfficientNetV2.ipynb

ใช้ Albumentations เพื่อสร้างการเพิ่มความหลากหลายในข้อมูลภาพสำหรับการ transform_train) และใช้การปรับขนาดภาพและทำ Normalization สำหรับข้อมูลทดสอบ

```
import albumentations as A
from albumentations.pytorch import ToTensorV2

transform_train = transforms.Compose([
    transforms.Resize((230,230)),
    transforms.RandomRotation(30),
    transforms.RandomCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.507, 0.487, 0.441], std=[0.267, 0.256, 0.276]) #normalize imagenet pretrain
])

transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.507, 0.487, 0.441], std=[0.267, 0.256, 0.276])
])

batch_size = 32
```

สร้างคลาส AnimalDataset สำหรับโมเดล PyTorch โดยใช้ Albumentations transforms สำหรับข้อมูลฝึก transform_train และ Albumentations transforms ที่ทำงานเฉพาะกับ PyTorch สำหรับข้อมูลทดสอบ

```
class AnimalDataset(Dataset):
    def __init__(self,
                 img_dir,
                 transforms=None):
        super().__init__()
        label_image = ['butterfly','cat','chicken','cow','dog','elephant','horse','sheep','spider','squirrel']
        self.input_dataset = list()
        label_num = 0
        for label in label_image:
            files = next(os.walk(os.path.join(img_dir,label)))
            for image_name in files:
                input = [os.path.join(img_dir,label,image_name),label_num] # [image_path, label_num]
                self.input_dataset.append(input)
                label_num += 1

        self.transforms = transforms

    def __len__(self):
        return len(self.input_dataset)

    def __getitem__(self, idx):
        img = Image.open(self.input_dataset[idx][0]).convert('RGB')
        x = self.transforms(img)
        y = self.input_dataset[idx][1]
        return x,y

trainset = AnimalDataset('./Dataset_animal2/train',transform_train)
valset = AnimalDataset('./Dataset_animal2/val',transform)
testset = AnimalDataset('./Dataset_animal2/test',transform)

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True)
valloader = torch.utils.data.DataLoader(valset, batch_size=batch_size, shuffle=True)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=True)

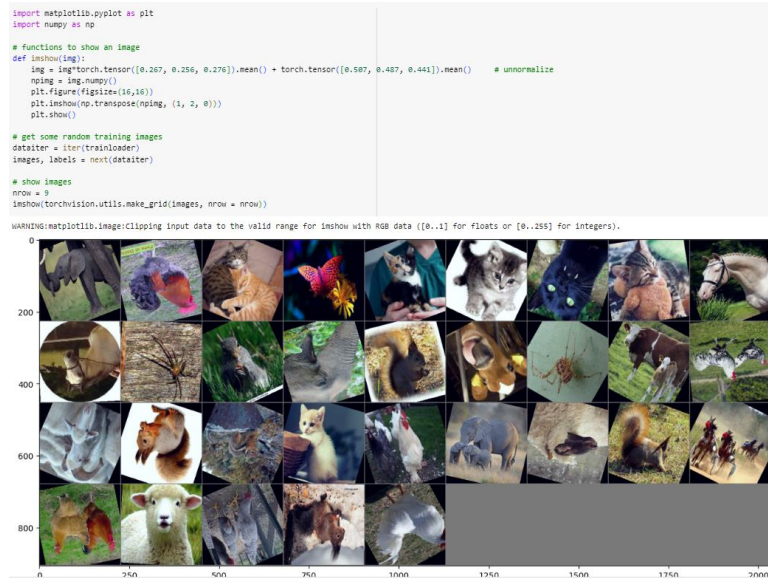
#classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

ใช้เพื่อดูขนาด (จำนวนตัวอย่าง) ของ Dataset สำหรับการฝึก (trainset), Dataset สำหรับ Validation และ Dataset สำหรับการทดสอบ (testset)

```
trainset.__len__(), valset.__len__(), testset.__len__()

(1400, 300, 300)
```


`def imshow(img):` เป็นการแสดงรูปภาพ ส่วน `iter(trainloader)` คือการสุ่มรูปภาพมาจาก `trainloader` โดยกำหนดให้ 1 แถวแสดง 9 ช่อง



พิมพ์ Label ของรูปภาพที่ถูกดึงมาจาก `trainloader` ในรูปแบบของกริด โดยใช้รูปเพื่อแสดง Label ของแต่ละรูปภาพในแต่ละแถวของกริด

```
# print labels
for i in range(batch_size//nrow + 1 if batch_size % nrow else 0):
    print(' '.join(f'{labels[i*nrow+j]:<3}' for j in range(min(batch_size - i*nrow, nrow))))
```

```
5  2  1  0  1  1  1  1  6
6  8  9  5  9  3  8  3  2
7  9  9  1  2  5  7  9  6
2  7  2  9  2
```

ใช้โมเดล EfficientNetV2 จาก `torchvision` และ โหลดน้ำหนักที่ถูกฝึกจาก ImageNet เพื่อใช้ในการจำแนกประเภท

```
import torch.nn as nn
import torch.nn.functional as F

pretrain_weight = torchvision.models.EfficientNet_V2_S_Weights.IMAGENET1K_V1
net = torchvision.models.efficientnet_v2_s(weights = pretrain_weight)
net.classifier[1] = nn.Linear(1280, 10)
net = net.to(device)
```

Downloading: "https://download.pytorch.org/models/efficientnet_v2_s-dd5fe13b.pth" to /root/.cache/torch/hub/checkpoints/efficientnet_v2_s-dd5fe13b.pth
100% 82.7M/82.7M [00:00<00:00, 165MB/s]

ใช้ไลบรารี `torchsummary` เพื่อแสดงสรุปข้อมูลโครงสร้างของโมเดล ด้วยขนาดของข้อมูลนำเข้าเป็น (3, 224, 224) และ batch size เท่ากับ 64

```
from torchsummary import summary
summary(net, (3, 224, 224), batch_size = 64)
```


Layer (type)	Output Shape	Param #		
Conv2d-1	[64, 24, 112, 112]	648	Conv2d-525	[64, 256, 7, 7]
BatchNorm2d-2	[64, 24, 112, 112]	48	BatchNorm2d-526	[64, 256, 7, 7]
SILU-3	[64, 24, 112, 112]	0	StochasticDepth-527	[64, 256, 7, 7]
Conv2d-4	[64, 24, 112, 112]	5,184	MBConv-528	[64, 256, 7, 7]
BatchNorm2d-5	[64, 24, 112, 112]	48	Conv2d-529	[64, 1536, 7, 7]
SILU-6	[64, 24, 112, 112]	0	BatchNorm2d-530	[64, 1536, 7, 7]
StochasticDepth-7	[64, 24, 112, 112]	0	SILU-531	[64, 1536, 7, 7]
FusedMBConv-8	[64, 24, 112, 112]	0	Conv2d-532	[64, 1536, 7, 7]
Conv2d-9	[64, 24, 112, 112]	5,184	BatchNorm2d-533	[64, 1536, 7, 7]
BatchNorm2d-10	[64, 24, 112, 112]	48	SILU-534	[64, 1536, 7, 7]
SILU-11	[64, 24, 112, 112]	0	AdaptiveAvgPool2d-535	[64, 1536, 1, 1]
StochasticDepth-12	[64, 24, 112, 112]	0	Conv2d-536	[64, 64, 1, 1]
FusedMBConv-13	[64, 24, 112, 112]	0	SILU-537	[64, 64, 1, 1]
Conv2d-14	[64, 96, 56, 56]	20,736	Conv2d-538	[64, 1536, 1, 1]
BatchNorm2d-15	[64, 96, 56, 56]	192	Sigmoid-539	[64, 1536, 1, 1]
SILU-16	[64, 96, 56, 56]	0	SqueezeExcitation-540	[64, 1536, 7, 7]
Conv2d-17	[64, 48, 56, 56]	4,608	Conv2d-541	[64, 256, 7, 7]
BatchNorm2d-18	[64, 48, 56, 56]	96	BatchNorm2d-542	[64, 256, 7, 7]
FusedMBConv-19	[64, 48, 56, 56]	0	StochasticDepth-543	[64, 256, 7, 7]
Conv2d-20	[64, 192, 56, 56]	82,944	MBConv-544	[64, 256, 7, 7]
BatchNorm2d-21	[64, 192, 56, 56]	384	Conv2d-545	[64, 1280, 7, 7]
SILU-22	[64, 192, 56, 56]	0	BatchNorm2d-546	[64, 1280, 7, 7]
Conv2d-23	[64, 48, 56, 56]	9,216	SILU-547	[64, 1280, 7, 7]
BatchNorm2d-24	[64, 48, 56, 56]	96	AdaptiveAvgPool2d-548	[64, 1280, 1, 1]
StochasticDepth-25	[64, 48, 56, 56]	0	Dropout-549	[64, 1280]
FusedMBConv-26	[64, 48, 56, 56]	0	Linear-550	[64, 10]
Conv2d-27	[64, 192, 56, 56]	82,944		
BatchNorm2d-28	[64, 192, 56, 56]	384		
SILU-29	[64, 192, 56, 56]	0		
Conv2d-30	[64, 48, 56, 56]	9,216		
BatchNorm2d-31	[64, 48, 56, 56]	96		
StochasticDepth-32	[64, 48, 56, 56]	0		
FusedMBConv-33	[64, 48, 56, 56]	0		
Conv2d-34	[64, 192, 56, 56]	82,944		
BatchNorm2d-35	[64, 192, 56, 56]	384		
SILU-36	[64, 192, 56, 56]	0		
Conv2d-37	[64, 48, 56, 56]	9,216		
BatchNorm2d-38	[64, 48, 56, 56]	96		

กำหนดฟังก์ชันความสูญเสียแบบ CrossEntropyLoss ตัวตั้งค่าการอัปเดตพารามิเตอร์ของโมเดลด้วยวิธีการ SGD โดยกำหนด learning rate เป็น 0.02 และ momentum เป็น 0.9. นอกจากนี้ยังกำหนด Scheduler ของ learning rate ด้วย StepLR ซึ่งลด learning rate ทุกระยะเวลาที่กำหนด (step_size=7) โดยการคูณ learning rate

```
import torch.optim as optim
from torch.optim import lr_scheduler

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.02, momentum=0.9)
scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.5)
```

ทำการฝึกโมเดล EfficientNetV2 ด้วยข้อมูลฝึกและตรวจสอบประสิทธิภาพในขณะที่บันทึกค่า loss, accuracy, และ f1-score ในแต่ละ epoch. โดยมีการใช้ CrossEntropyLoss เป็นฟังก์ชันความสูญเสีย, SGD เป็นตัวตั้งค่าการอัปเดตพารามิเตอร์, และ Scheduler ของ learning rate เพื่อลด learning rate

```
from sklearn.metrics import classification_report
from tqdm.notebook import tqdm

epochs = 20

history_train = {'loss': np.zeros(epochs), 'acc': np.zeros(epochs), 'f1-score': np.zeros(epochs)}
history_val = {'loss': np.zeros(epochs), 'acc': np.zeros(epochs), 'f1-score': np.zeros(epochs)}
min_val_loss = 1e10
PATH = './efficientnetv2.pth'

for epoch in range(epochs): # loop over the dataset multiple times

    print(f'epoch {epoch + 1} \ntraining ...')
    net.train()
    y_predict = list()
    y_labels = list()
    training_loss = 0.0
    n = 0
    with torch.set_grad_enabled(True):
        for data in tqdm(trainloader):
            # get the inputs; data is a list of [inputs, labels]
            inputs, labels = data
            inputs = inputs.to(device)
            labels = labels.to(device)

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # aggregate statistics
            training_loss += loss.item()
            n += 1

    y_labels += list(labels.cpu().numpy())
    y_predict += list(outputs.argmax(dim=-1).cpu().numpy())
    scheduler.step()

    # print statistics
    report = classification_report(y_labels, y_predict, digits = 4, output_dict = True)
    acc = report['accuracy']
    f1 = report['f1-score']
    support = report['support']
    training_loss /= n
    print(f'training loss: {training_loss:.4}, acc: {acc*100:.4}%, f1-score: {f1*100:.4}%, support: {support}')
    history_train['loss'][epoch] = training_loss
    history_train['acc'][epoch] = acc
    history_train['f1-score'][epoch] = f1

    print('validating ...')
    net.eval()
    optimizer.zero_grad()
    y_predict = list()
    y_labels = list()
    validation_loss = 0.0
    n = 0
    with torch.no_grad():
        for data in tqdm(valloader):
            inputs, labels = data
            inputs = inputs.to(device)
            labels = labels.to(device)

            outputs = net(inputs)
            loss = criterion(outputs, labels)
            validation_loss += loss.item()

    y_labels += list(labels.cpu().numpy())
    y_predict += list(outputs.argmax(dim=-1).cpu().numpy())
    n += 1

    # print statistics
    report = classification_report(y_labels, y_predict, digits = 4, output_dict = True)
    acc = report['accuracy']
    f1 = report['f1-score']
    support = report['support']
    validation_loss /= n
    print(f'validation loss: {validation_loss:.4}, acc: {acc*100:.4}%, f1-score: {f1*100:.4}%, support: {support}')
    history_val['loss'][epoch] = validation_loss
    history_val['acc'][epoch] = acc
    history_val['f1-score'][epoch] = f1

    # save min validation loss
    if validation_loss < min_val_loss:
        torch.save(net.state_dict(), PATH)
        min_val_loss = validation_loss

    print(f'Finished Training')
```

```

epoch 1
Training ...
100% ██████████ 44/44 [00:24<00:00, 2.44it/s]
training loss: 0.9199, acc: 72.71%, f1-score: 73.12%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.87it/s]
validation loss: 0.7861, acc: 79.0%, f1-score: 79.29%, support: 300
epoch 2
Training ...
100% ██████████ 44/44 [00:18<00:00, 2.50it/s]
training loss: 0.6323, acc: 80.5%, f1-score: 80.4%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.87it/s]
validation loss: 0.436, acc: 84.67%, f1-score: 84.7%, support: 300
epoch 3
Training ...
100% ██████████ 44/44 [00:18<00:00, 2.53it/s]
training loss: 0.3498, acc: 89.07%, f1-score: 89.09%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 4.82it/s]
validation loss: 0.4421, acc: 86.0%, f1-score: 86.12%, support: 300
epoch 4
Training ...
100% ██████████ 44/44 [00:18<00:00, 2.53it/s]
training loss: 0.3235, acc: 89.93%, f1-score: 89.91%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.90it/s]
validation loss: 0.2992, acc: 89.0%, f1-score: 89.05%, support: 300
epoch 5
Training ...
100% ██████████ 44/44 [00:19<00:00, 2.53it/s]
training loss: 0.2004, acc: 93.57%, f1-score: 93.58%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.82it/s]
validation loss: 0.369, acc: 89.33%, f1-score: 89.29%, support: 300
epoch 6
Training ...

```

```

validation loss: 0.2485, acc: 95.0%, f1-score: 94.96%, support: 300
epoch 16
Training ...
100% ██████████ 44/44 [00:19<00:00, 2.34it/s]
training loss: 0.0124, acc: 99.64%, f1-score: 99.64%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.61it/s]
validation loss: 0.2341, acc: 95.0%, f1-score: 94.98%, support: 300
epoch 17
Training ...
100% ██████████ 44/44 [00:19<00:00, 2.47it/s]
training loss: 0.009672, acc: 99.71%, f1-score: 99.71%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.04it/s]
validation loss: 0.2359, acc: 95.33%, f1-score: 95.31%, support: 300
epoch 18
Training ...
100% ██████████ 44/44 [00:19<00:00, 2.45it/s]
training loss: 0.01126, acc: 99.71%, f1-score: 99.71%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.53it/s]
validation loss: 0.234, acc: 95.0%, f1-score: 95.01%, support: 300
epoch 19
Training ...
100% ██████████ 44/44 [00:19<00:00, 2.43it/s]
training loss: 0.0104, acc: 99.71%, f1-score: 99.71%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.69it/s]
validation loss: 0.231, acc: 94.33%, f1-score: 94.33%, support: 300
epoch 20
Training ...
100% ██████████ 44/44 [00:19<00:00, 2.39it/s]
training loss: 0.01399, acc: 99.36%, f1-score: 99.36%, support: 1400
validating ...
100% ██████████ 10/10 [00:01<00:00, 5.54it/s]
validation loss: 0.2241, acc: 94.67%, f1-score: 94.65%, support: 300
Finished Training

```

min_val_loss เก็บค่า validation loss ที่ต่ำที่สุดที่พบในระหว่างการฝึกโมเดล

min_val_loss


0.22409853991121054

ใช้ Matplotlib เพื่อสร้างกราฟที่แสดงค่า loss, accuracy, และ f1-score ของโมเดลในแต่ละ epoch จาก training data และ test data

```

fig, axs = plt.subplots(3, figsize= (6,10))
# loss
axs[0].plot(history_train['loss'], label = 'training')
axs[0].plot(history_val['loss'], label = 'validation')
axs[0].set_title("loss")
axs[0].legend()
# acc
axs[1].plot(history_train['acc'], label = 'training')
axs[1].plot(history_val['acc'], label = 'validation')
axs[1].set_title("acc")
axs[1].legend()
# f1-score
axs[2].plot(history_train['f1-score'], label = 'training')
axs[2].plot(history_val['f1-score'], label = 'validation')
axs[2].set_title("f1-score")
axs[2].legend()
plt.show()

```



ใช้ PyTorch เพื่อโหลดน้ำหนัก (weights) ของโมเดล EfficientNetV2 ที่ถูกบันทึกไว้ในไฟล์ที่กำหนด (PATH) เพื่อนำมาใช้ในการทดสอบ

```
PATH = './Animal10-efficientnetV2s.pth'
```

```
net.load_state_dict(torch.load(PATH))
```

```
<All keys matched successfully>
```

ทำการทดสอบโมเดล EfficientNetV2 ด้วยข้อมูลทดสอบและแสดงผลพร้อมด้วยตัวช่วยที่เกี่ยวข้องกับ Scikit-Learn

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

print('testing ...')
y_predict = list()
y_labels = list()
test_loss = 0.0
n = 0
with torch.no_grad():
    for data in tqdm(testloader):
        net.eval()
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)

        outputs = net(inputs)
        loss = criterion(outputs, labels)
        test_loss += loss.item()

    y_labels += list(labels.cpu().numpy())
    y_predict += list(outputs.argmax(dim=1).cpu().numpy())
    n+=1

# print statistics
test_loss /= n
print(f"testing loss: {test_loss:.4}" )

report = classification_report(y_labels, y_predict, digits = 4)
M = confusion_matrix(y_labels, y_predict)
print(report)
disp = ConfusionMatrixDisplay(confusion_matrix=M)
```

```
testing ...
100% ██████████ 10/10 [00:01<00:00, 5.56it/s]
testing loss: 0.2259
```

	precision	recall	f1-score	support
0	0.8788	0.9667	0.9206	30
1	1.0000	0.9000	0.9474	30
2	0.9333	0.9333	0.9333	30
3	0.9355	0.9667	0.9508	30
4	0.8824	1.0000	0.9375	30
5	1.0000	0.9667	0.9831	30
6	0.9355	0.9667	0.9508	30
7	1.0000	0.9000	0.9474	30
8	0.9643	0.9000	0.9310	30
9	1.0000	1.0000	1.0000	30
accuracy			0.9500	300
macro avg	0.9530	0.9500	0.9502	300
weighted avg	0.9530	0.9500	0.9502	300

ใช้ Matplotlib เพื่อแสดงแผนภูมิ confusion matrix ที่ถูกสร้างขึ้นโดย ConfusionMatrixDisplay

```
disp.plot()  
plt.show()
```

